



MANGOBOOST

AMDSOW InferenceX documentation set

**Start here, choose the right guide, then run
the benchmark**

AMDSOW InferenceX DeepSeek-R1-0528 FP8 MI300X vLLM PD-
disaggregation validation package

Confidential. Shared with AMD under NDA. | contact@mangoboost.io

Document control

Field	Value
Document	AMDSOW InferenceX documentation set
Version	v1.0.0
Last updated	2026-06-25
Owner	MangoBoost AMDSOW delivery team
Contact	contact@mangoboost.io
Audience	AMD/customer operators, solution engineers, maintainers, and package reviewers
Status	Generated A4 PDF from maintained AMDSOW InferenceX Markdown documentation
Source format	Maintained Markdown rendered through HTML/CSS with Python-Markdown and WeasyPrint.

AMDSOW InferenceX onboarding

Start here if you are seeing this repository for the first time.

In plain terms: this documentation explains how to run the AMDSOW DeepSeek-R1-0528 FP8 MI300X vLLM prefill/decode benchmark through InferenceX. GitHub Actions starts the run, a self-hosted runner submits a Slurm job, Slurm allocates MI300X compute nodes, Docker starts vLLM workers, and the benchmark writes JSON results.

Five-minute checklist

1. Read `AMDSOW_INFERENCEX_USER_GUIDE.md` sections 1-4 before dispatching anything.
2. Confirm your local shell has `.venv/bin/python` for matrix validation and `gh` for GitHub Actions dispatch.
3. Validate the default row locally: `8k1k` , concurrency `256` , config key `dsr1-fp8-mi300x-vllm-disagg` .
4. Confirm the shared `mi300x-disagg` runner and target Slurm cluster are idle using the site-approved status source. `squeue` is useful only from a Slurm-capable host and may not show every operator's jobs.
5. Use `AMDSOW_INFERENCEX_PROJECT_DESCRIPTION_DETAIL.md` before changing YAML, Slurm, Docker, or vLLM flags.

Documents and PDFs

Need	Markdown	PDF
Start here and choose the right document	<code>AMDSOW_INFERENCEX_ONBOARDING.md</code>	<code>00-amdsow-inferencex-documentation-set.pdf</code>
Run the benchmark, copy commands, monitor results, troubleshoot an operator run	<code>AMDSOW_INFERENCEX_USER_GUIDE.md</code>	<code>01-amdsow-inferencex-user-guide.pdf</code>
Understand the project wiring, config semantics, script flow, code excerpts, and safe edit rules	<code>AMDSOW_INFERENCEX_PROJECT_DESCRIPTION_DETAIL.md</code>	<code>02-amdsow-inferencex-project-detail.pdf</code>

Use the user guide first if you are operating a benchmark. Use the project detail document only when you need internals or are changing configs/scripts.

Source-of-truth rule

The source of truth for the AMDSOW InferenceX benchmark matrix is `.github/configs/amd-master.yaml` , key `dsr1-fp8-mi300x-vllm-disagg` . Local wrappers and validation scripts are convenience tools; do not treat them as the canonical config.